


Because of
you,

AGILE,

I Need to
Learn ---->




2110531 - Data Science and Data Engineering Tools

Ops Stars

Asst.Prof. Natawut Nupairoj, Ph.D.
Department of Computer Engineering
Chulalongkorn University
natawut.n@chula.ac.th


Agile Development and the Real World Problems




Solution: DevOps

- A set of practices that combines software development (Dev) and IT operations (Ops)
- Aim to shorten SDLC with continuous delivery and high software quality
 - Tackle deployment architecture at early stages
 - Involve everyone in SDLC in a continuous fashion
 - Automate all necessary steps as much as and as early as possible
- Software-defined system (Infrastructure as Code) plays a very important role to ensure smooth integration and delivery

The Infinite Lifecycle




DevOps Automation




Source: <https://www.bmc.com/blogs/continuous-delivery-continuous-deployment-continuous-integration-whats-difference/>

With ML and Data Science, these cycles become more complicated




DataOps and MLOps




Orchestrating Data Pipeline for DataOps

Ops Stars




DataOps and MLOps



Apache Airflow

- Open-source workflow management platform for data engineering pipelines, started by Airbnb in 2014
- Designed under the principle of "configuration as code"
 - Uses directed acyclic graphs (DAGs) to manage workflow orchestration
 - Tasks and dependencies are defined in Python (possible for just one file)
 - DAG can be started by defined calendar schedule or triggered by external event
 - Airflow manages the scheduling and execution
- Provide managing and monitoring facility with web-based UI

Flexibility of Data Pipelines as Code




Complex pipeline can be implemented using operators, provider packages, and external services — with minimal code writing

Provider Packages


Providers packages include integrations with third party projects. They are updated independently of the Apache Airflow core. [Read the documentation >>](#)

- [Airbyte](#)
- [Alibaba](#)
- [Amazon](#)
- [Apache Beam](#)
- [Apache Cassandra](#)
- [Apache Drill](#)
- [Apache Druid](#)
- [Apache HDFS](#)
- [Apache Hive](#)
- [Apache Kylin](#)
- [Apache Livy](#)
- [Apache Pig](#)
- [Apache Pinot](#)
- [Apache Spark](#)
- [Apache Sqoop](#)
- [Asana](#)
- [Celery](#)
- [IBM Cloudant](#)
- [Kubernetes](#)
- [Databricks](#)
- [Datadog](#)
- [DBT cloud](#)
- [Dingding](#)
- [Discord](#)
- [Docker](#)
- [Elasticsearch](#)
- [Exasol](#)
- [Facebook](#)
- [File Transfer Protocol \(FTP\)](#)
- [Github](#)
- [Google](#)
- [gRPC](#)
- [Hashicorp](#)
- [Hypertext Transfer Protocol \(HTTP\)](#)
- [Influx DB](#)
- [Internet Message Access Protocol \(IMAP\)](#)
- [Java Database Connectivity \(JDBC\)](#)
- [Jenkins](#)
- [Jira](#)
- [Microsoft Azure](#)
- [Microsoft PowerShell Remoting Protocol \(PSRP\)](#)
- [Microsoft SQL Server \(MSSQL\)](#)
- [Windows Remote Management \(WinRM\)](#)
- [MongoDB](#)
- [MySQL](#)
- [Neo4J](#)
- [ODBC](#)
- [OpenFaaS](#)
- [Opsgenie](#)
- [Oracle](#)
- [Pagerduty](#)
- [Papermill](#)
- [Plexus](#)
- [PostgreSQL](#)
- [Presto](#)
- [Qubole](#)
- [Redis](#)
- [Salesforce](#)
- [Samba](#)
- [Segment](#)
- [Sendgrid](#)
- [SFTP](#)
- [Singularity](#)
- [Slack](#)
- [Snowflake](#)
- [SQLite](#)
- [SSH](#)
- [Tableau](#)
- [Telegram](#)
- [Trino](#)
- [Vertica](#)
- [Yandex](#)
- [Zendesk](#)

Airflow Architecture



Airflow Components



- **DAG** (directed acyclic graph) is a collection of tasks in pipeline
- **Task** is the basic unit of execution, an instantiated operator object with unique task id, and can be arranged into DAG to define dependencies among tasks
- **Operator** is a predefined task template being used in many situations e.g. BashOperator, PythonOperator, EmailOperator, Sensors, TaskFlow, and others (database operators, cloud storage operators, etc.)
- Others: hooks, coms, trigger, pool, variables, etc.

Example DAG

```
from datetime import datetime

from airflow import DAG
from airflow.decorators import task
from airflow.operators.bash import BashOperator

# A DAG represents a workflow, a collection of tasks
with DAG(dag_id="demo", start_date=datetime(2022, 1, 1), schedule="0 0 * * *") as dag:
    # Tasks are represented as operators
    hello = BashOperator(task_id="hello", bash_command="echo hello")

    @task()
    def airflow():
        print("airflow")

    # Set dependencies between tasks
    hello >> airflow()
```

How to run airflow

- Two recommended approaches, standalone and docker

```
pip install apache-airflow
```

```
airflow db init
```

```
airflow users create --username admin --firstname Natawut --  
lastname Nupairoj --role Admin --email natawut.n@chula.ac.th
```

```
airflow standalone
```

Airflow UI

Airflow

DAGs

Security

Browse

Admin

Docs

09:32 UTC

NN

Do not use **SQLite** as metadata DB in production – it should only be used for dev/testing. We recommend using Postgres or MySQL. [Click here](#) for more information.

Do not use **SequentialExecutor** in production. [Click here](#) for more information.

DAGs

All 34

Active 0

Paused 34

Filter DAGs by tag

Search DAGs

DAG	Owner	Runs	Schedule	Last Run	Next Run	Recent Tasks
concurrent_dag	airflow	1 2	1 day, 0:00:00	2022-03-28, 08:33:25	2022-03-28, 00:00:00	
cp_webcheck	airflow	3	1 day, 0:00:00	2022-03-28, 08:59:01	2022-03-28, 00:00:00	
cp_webcheck_failed	airflow	2	1 day, 0:00:00	2022-03-28, 09:00:46	2022-03-28, 00:00:00	
example example2 example_bash_operator	airflow	2	0 0 ***	2015-01-02, 00:00:00	2022-03-27, 00:00:00	
example example_branch_datetime_operator_2	airflow		@daily		2022-03-27, 00:00:00	
example example_branch_dop_operator_v3	airflow		* / 1 * * * *		2022-03-28, 09:30:00	
example_branch_labels	airflow		@daily		2022-03-27, 00:00:00	

Example: Simple BashOperator

dsde_simplebash.py

```
from airflow.utils.dates import days_ago
from airflow import DAG
from airflow.operators.bash import BashOperator

# declare a DAG
with DAG(dag_id='dsde_simplebash', start_date=days_ago(1)):

    # declare a TASK
    echo = BashOperator(task_id='echo_template',
bash_command='echo "run_id = {{ run_id }} and ds = {{ ds }}"')
```


BashOperator

- BashOperator allows Airflow to run a task using bash script
- The exit code of the program/script will determine the result of the task execution
 - 0 - Success
 - 99 - Skipped
 - Something else - Failed
- Parameterizing bash_command is possible using Jinja templates (refer to <https://airflow.apache.org/docs/apache-airflow/stable/templates-ref.html> for more details)
- Airflow captures the last line of the output from the script in 'output'

Example: Task Dependency

dsde_simplebash2.py

```
from airflow.utils.dates import days_ago
from airflow import DAG
from airflow.operators.bash import BashOperator

# declare a DAG
with DAG(dag_id='dsde_simplebash_2', start_date=days_ago(1)):

    # declare 2 TASKS
    echo = BashOperator(task_id='echo_template', bash_command='echo "run_id
= {{ run_id }}" and ds = {{ ds }}"')
    echo2 = BashOperator(task_id='echo_template_2', bash_command='echo "[2]
run_id = {{ run_id }}" and ds = {{ ds }}"')

    # declare dependency
    echo >> echo2
```

Example: Concurrent Tasks

dsde_concurrent.py

```
from airflow.utils.dates import days_ago
from airflow import DAG

from airflow.operators.bash import BashOperator
from airflow.operators.dummy import DummyOperator

# 'with' enables DAG to become context managers; automatically assign new operators to
that DAG
with DAG('concurrent_dag', start_date=days_ago(1)) as dag:
    start = DummyOperator(task_id='start_task')
    ping = BashOperator(task_id='cp_check', bash_command='curl https://
www.cp.eng.chula.ac.th')
    ping2 = BashOperator(task_id='eng_check', bash_command='curl https://
www.eng.chula.ac.th')
    inform = BashOperator(task_id='inform_status', bash_command='echo "CP website still
works!"')

    # creating DAG dependencies can be a long flow or multiple short flows
    # start >> [ping, ping2] >> inform
    start >> [ping, ping2]
    ping >> inform
    ping2 >> inform
```

Airflow Common Operators

- BashOperator - executes a bash command
- PythonOperator - calls an arbitrary Python function
- EmailOperator - sends an email
- SimpleHttpOperator - sends an HTTP request
- MySqlOperator, SqliteOperator, PostgresOperator, MsSqlOperator, OracleOperator, JdbcOperator, etc. - executes a SQL command
- Sensor - an Operator that waits (polls) for a certain time, file, database row, S3 key, etc...
- Other specific operators: DockerOperator, HiveOperator, S3FileTransformOperator, PrestoToMySqlTransfer, SlackAPIOperator, etc.

PythonOperator / PythonVirtualenvOperator

- PythonOperator provides more flexibility of running python code e.g. perform data conversion, train a model
- Input parameters can be passed via op_kwargs/op_args and XCom
- Return values are captured in 'output' variable

Example: PythonOperator

dsde_python.py

```
from airflow.utils.dates import days_ago
from airflow import DAG


from airflow.operators.bash import BashOperator
from airflow.operators.dummy import DummyOperator
from airflow.operators.python import PythonOperator

def show_status(web1, web2):
    print('All websites are running!')
    print('web1 = ', web1)
    print('web2 = ', web2)
    print('--- DONE ---')

# 'with' enables DAG to become context managers; automatically assign new operators to that DAG
with DAG('python_operator', start_date=days_ago(1)) as dag:
    start = DummyOperator(task_id='start_task')
    ping = BashOperator(task_id='cp_check', bash_command='curl https://www.cp.eng.chula.ac.th')
    ping2 = BashOperator(task_id='eng_check', bash_command='curl https://www.eng.chula.ac.th')
    inform = PythonOperator(task_id='inform_status', python_callable=show_status,
op_args=[ping.output, ping2.output])

    # creating DAG dependencies can be a long flow or multiple short flows
    # start >> [ping, ping2] >> inform
    start >> [ping, ping2]
    ping >> inform
    ping2 >> inform
```

BranchPythonOperator



- Allows to follow a specific path according to a condition, evaluated in a python callable function
- Executes a Python function returning a list of task ids of the next tasks to execute

Example: PythonBranchOperator

dsde_branch.py

```
from airflow import DAG
from airflow.operators.python import BranchPythonOperator
from airflow.operators.dummy import DummyOperator
from datetime import datetime

def _choose_best_model(accuracy):
    if accuracy > 0.8:
        return 'accurate'
    return 'inaccurate'

with DAG('branch_oper', start_date=datetime(2021, 1, 1), catchup=False) as dag:
    choose_best_model = BranchPythonOperator(task_id='choose_best_model',
python_callable=_choose_best_model, op_args=[0.75])

    accurate = DummyOperator(task_id='accurate')
    inaccurate = DummyOperator(task_id='inaccurate')

    choose_best_model >> [accurate, inaccurate]
```

See more DAG examples in airflow

More about DAG

- DAG can be declared with the following parameters
 - `start_date` - when to start, can be exact date or relative e.g. `days_ago(1)`
 - `schedule_interval` - how often should this DAG run
 - `catchup` - should scheduler kick off a DAG Run for any data interval that has not been run since the last data interval
 - `default_args` - args to get passed on to each operator in this DAG, can override them on a per-task basis
- Airflow 2.0 introduces TaskFlow API with DAG and Task decorators

Sensors

- a special type of Operator that are designed to wait until something happens, and then succeed so their downstream tasks can run
- Can be time-based, or waiting for a file, or an external event
- Sensors have three different modes of running
 - poke (default): The Sensor takes up a worker slot for its entire runtime
 - reschedule: The Sensor takes up a worker slot only when it is checking, and sleeps for a set duration between checks
 - smart sensor: There is a single centralized version of this Sensor that batches all executions of it

Other Components


- Connectors - repository of configurations to connect to external systems
- Hooks - high-level interface (driver) to an external platform
- Variables - generic way to store and retrieve configured values (key-value pair), similar to environment variables
- XCom - operator cross-communication - let tasks exchange information
- Executors - mechanism how task instances get to run (default = LocalExecutors)
- Trigger - run DAG from CLI, REST API, Web UI, TriggerDagRunOperator

Use Cases


- ETL Pipelines that extract data from multiple sources and run Spark jobs or any other data transformations
- Collecting Sensor data and move to Data Lake/Data warehouse
- Training machine learning models
- Orchestrating automated testing
- Report generation

Orchestrating ML Pipeline

Ops Stars




DataOps and MLOps



MLOps Orchestration Requirements

- In addition to workflow control, orchestrating MLOps requires some unique requirements
 - Code + Data = combination of development and data pipeline
 - Model version control
 - Model performance tracking during training (and serving)

MLflow Components



Key Concepts in MLflow

- **A run** is a collection of parameters, metrics, labels, and artifacts related to the training process of a machine learning model
- **An experiment** is the basic unit of MLflow organization ; all MLflow runs belong to an experiment
- For each experiment, you can analyze and compare the results of different runs, and easily retrieve metadata artifacts for analysis using downstream tools

How to run MLflow

- Install the MLflow

```
pip install mlflow
```

- Install the tutorial code

```
git clone https://github.com/mlflow/mlflow
```

Harutaka Kawamura, 2 years ago | 4 authors (Matei Zaharia and others)

```
1 # The data set used in this example is from http://archive.ics.uci.edu/ml/datasets/Wine+Quality
2 # P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis.
3 # Modeling wine preferences by data mining from physicochemical properties. In Decision Support Systems, Elsevier, 47(4):547-
4
5 import os
6 import warnings
7 import sys
8
9 import pandas as pd
10 import numpy as np
11 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
12 from sklearn.model_selection import train_test_split
13 from sklearn.linear_model import ElasticNet
14 from urllib.parse import urlparse
15 import mlflow
16 import mlflow.sklearn
17
18 import logging
19
20 logging.basicConfig(level=logging.WARN)
21 logger = logging.getLogger(__name__)
22
23
24 def eval_metrics(actual, pred):
25     rmse = np.sqrt(mean_squared_error(actual, pred))
26     mae = mean_absolute_error(actual, pred)
27     r2 = r2_score(actual, pred)
28     return rmse, mae, r2
29
30
31 if __name__ == "__main__":
```

```
58 with mlflow.start_run():
59     lr = ElasticNet(alpha=alpha, l1_ratio=l1_ratio, random_state=42)
60     lr.fit(train_x, train_y)
61
62     predicted_qualities = lr.predict(test_x)
63
64     (rmse, mae, r2) = eval_metrics(test_y, predicted_qualities)
65
66     print("Elasticnet model (alpha=%f, l1_ratio=%f):" % (alpha, l1_ratio))
67     print("  RMSE: %s" % rmse)
68     print("  MAE: %s" % mae)
69     print("  R2: %s" % r2)
70
71     mlflow.log_param("alpha", alpha)
72     mlflow.log_param("l1_ratio", l1_ratio)
73     mlflow.log_metric("rmse", rmse)
74     mlflow.log_metric("r2", r2)
75     mlflow.log_metric("mae", mae)
76
77     tracking_url_type_store = urlparse(mlflow.get_tracking_uri()).scheme
78
79     # Model registry does not work with file store
80     if tracking_url_type_store != "file":
81
82         # Register the model
83         # There are other ways to use the Model Registry, which depends on the use case,
84         # please refer to the doc for more information:
85         # https://mlflow.org/docs/latest/model-registry.html#api-workflow
86         mlflow.sklearn.log_model(lr, "model", registered_model_name="ElasticnetWineModel")
87     else:
88         mlflow.sklearn.log_model(lr, "model")
```

Running Model and View Results

- At mlflow tutorial code repo:

```
cd mlflow/examples
python sklearn_elasticnet_wine/train.py
python sklearn_elasticnet_wine/train.py [alpha] [l1_ratio]
```

- View results (run at mlflow/examples) and use your browser to <http://127.0.0.1:5000>

```
cd mlflow/examples
mlflow ui
```


mlflow1.30.0

ExperimentsModels

GitHubDocs

Experiments

Search Experiments

Default

Default

Share

Track machine learning training runs in experiments. [Learn more](#)

Experiment ID: 0

Description

Edit

Refresh

Compare

Delete

Download CSV

Created

All time

Columns

Only show differences

metrics.rmse < 1 and params.model = "tree"

Search

Filter

Clear

Showing 5 matching runs

								Metrics			Parameters	
	Created	Duration	Run Name	User	Source	Version	Models	mae	r2	rmse	alpha	l1_ratio
	4 minutes ago	2.5s	selective-o...	natawut	train.py	fe60af	sklearn	0.627	0.109	0.793	0.5	0.5
	16 minutes ago	2.4s	bedecked-...	natawut	train.py	fe60af	sklearn	0.673	0.017	0.833	1.0	1.0
	16 minutes ago	2.5s	rare-calf-8...	natawut	train.py	fe60af	sklearn	0.627	0.109	0.793	0.5	0.5
	31 minutes ago	2.6s	redolent-cr...	natawut	train.py	fe60af	sklearn	0.673	0.017	0.833	1.0	1.0
	31 minutes ago	2.5s	bedecked-...	natawut	train.py	fe60af	sklearn	0.627	0.109	0.793	0.5	0.5

Load more

Packaging and Serving Model

- Model can be packaged for running with its own python environment (for other to test or run on cloud)

```
mlflow run sklearn_elasticnet_wine -P alpha=0.42
```

- Model can be deployed as a standalone service by selecting the run that we want to use

```
mlflow models serve -m ./mlruns/0/12334a9519a94de99a3363341b69c97c/  
artifacts/model -p 1234
```

```
curl -X POST -H "Content-Type:application/json; format=pandas-split" --  
data '{"columns":["alcohol", "chlorides", "citric acid", "density",  
"fixed acidity", "free sulfur dioxide", "pH", "residual sugar",  
"sulphates", "total sulfur dioxide", "volatile acidity"],"data":[[12.8,  
0.029, 0.48, 0.98, 6.2, 29, 3.33, 1.2, 0.39, 75, 0.66]]}' http://  
127.0.0.1:1234/invocations
```

mlflow1.30.0

ExperimentsModels

GitHubDocs

Default > selective-owl-216

selective-owl-216

Run ID: 12334a9519a94de99a3363341b69c97c

Date: 2022-10-31 20:48:19

Source: train.py

Git Commit: fe60afd98c8a783ed3ced24150f2784f980ce790

User: natawut

Duration: 2.5s

Status: FINISHED

Lifecycle Stage: active

> Description Edit

> Parameters (2)

> Metrics (3)

> Tags

> Artifacts

model

MLmodelconda.yamlmodel.pklpython_env.yamlrequirements.txt

Full Path:./mlruns/0/12334a9519a94de99a3363341b69c97c/artifacts/model

Register Model

MLflow Model

The code snippets below demonstrate how to make predictions using the logged model. You can also register it to the model registry to version control

Model schema

Input and output schema for your model. Learn more

Name	Type
No schema. See MLflow docs for how to include input and output schema with your model.	

Make Predictions

Predict on a Spark DataFrame:

```
import mlflow
from pyspark.sql.functions import struct, col
logged_model = 'runs:/12334a9519a94de99a3363341b69c97c/model'

# Load model as a Spark UDF. Override result_type if the model does not return double values.
loaded_model = mlflow.pyfunc.spark_udf(spark, model_uri=logged_model, result_type='double')

# Predict on a Spark DataFrame.
df.withColumn('predictions', loaded_model(struct(*map(col, df.columns))))
```

Predict on a Pandas DataFrame:


Provisioning ML Model

Ops Stars



Provisioning Machine Learning Model

- Encapsulating your model as a microservice is a popular approach to deploy machine learning model
- The model becomes a black box and can be invoked using REST protocol
- Using a microframework such as FastAPI can simplify microservice development



Microframework

- Minimalistic web application frameworks designed for Microservices and API development
- Facilitate receiving an HTTP request, routing the HTTP request to the appropriate controller, dispatching the controller, and returning an HTTP response
- Designed for building the APIs for another service or application; therefore, lacking most of the functionality expected in full-stack framework

REST - Heart and Soul of Microservices

- REpresentational State Transfer
- Lightweight and simple
 - Only simple XML or JSON on HTTP
 - Utilize HTTP methods instead of using “envelope” style
 - URI as method identification
 - Emulate CRUD operations

REST over HTTP



FastAPI

- Web framework for developing RESTful APIs in Python based on Pydantic and starlette
- Provide several facilities for REST development
 - Data validation / serialization / deserialization
 - Auto-generate OpenAPI documents
 - Asynchronous programming
 - Dependency Injection
 - Support GraphQL for complex data model
 - Run with Uvicorn and Unicorn

Install Fastapi

- Very simple installation process with pip
- Make sure that you install both fastapi and unicorn

```
pip install fastapi uvicorn
```

Example: Simple Service

```
# simple fastapi example
```

```
import uvicorn
from fastapi import FastAPI
```

```
app = FastAPI()
```



```
@app.get("/")
async def root():
    return {"message": "My first FastAPI service"}
```

```
if __name__ == "__main__":
    uvicorn.run('simple:app', host="0.0.0.0", port=8000, reload=True)
```


Testing Our First Service

- Test the service with your browser at <http://localhost:8000>
- More information at <http://localhost:8000/docs>
- You can also try <http://localhost:8000/redoc>
- Beware of browser cache
- You can also use other tools: curl, postman, etc.

```
{"message": "My first FastAPI service"}
```



Example: Path Parameter with Type Validation

```
# path parameter example
```

```
import uvicorn
from fastapi import FastAPI
```

```
app = FastAPI()
```

```
@app.get('/items/{item_id}')
async def read_item(item_id: int):
    return { 'item_id': item_id }
```

```
if __name__ == "__main__":
    uvicorn.run('path_param:app', host="0.0.0.0", port=8000, reload=True)
```

Test this service at: <http://localhost:8000/items/1> and api document at: <http://localhost:8000/docs>

REST CRUD Operations

HTTP	CRUD	Decorator
GET	Read	<code>@app.get ()</code>
POST	Create, Update, Delete	<code>@app.post ()</code>
PUT	Create, Overwrite / Replace	<code>@app.put ()</code>
DELETE	Delete	<code>@app.delete ()</code>

Note

- FastAPI doesn't enforce any specific meaning
- Other operations (options, head, patch, trace) are also supported

POST Operation

- POST operation allows passing complex input data in JSON via request body
- Declaring data type with pydantic, FastAPI will
 - Read the body of the request as JSON
 - Convert the corresponding types if needed
 - Validate the data and return a clear error if it is invalid
 - Put the received data in the input parameter of the function

Example: POST for Complex Input Parameter

post example

```
import uvicorn
from fastapi import FastAPI
```

```
from typing import Optional
from pydantic import BaseModel
```

```
class Item(BaseModel):
    name: str
    description: Optional[str] = None
    price: float
    tax: Optional[float] = None
```

```
app = FastAPI()
```

```
@app.post("/items/")
async def create_item(item: Item):
    return item
```

```
@app.get('/items/{item_id}')
async def read_item(item_id: int):
    return { 'item_id': item_id }
```

```
if __name__ == "__main__":
    uvicorn.run('post:app', host="0.0.0.0", port=8000, reload=True)
```

Example: Using Path Parameter and Request Body

```
# put example
```

```
import uvicorn
from fastapi import FastAPI
```

```
from typing import Optional
from pydantic import BaseModel
```


```
class Item(BaseModel):
    name: str
    description: Optional[str] = None
    price: float
    tax: Optional[float] = None
```

```
app = FastAPI()
```

```
@app.put("/items/{item_id}")
async def update_item(item_id: int, item: Item):
    return {"item_id": item_id, **item.dict()}
```

```
if __name__ == "__main__":
    uvicorn.run('put:app', host="0.0.0.0", port=8000, reload=True)
```

DataOps and MLOps Full Loop



Conclusion

- Data pipeline can be very complex as it requires task orchestration across multiple platforms
- Airflow provides sophisticated data pipeline orchestration system with configuration as code and other comprehensive facilities
- Airflow is good for orchestrating data pipeline, ETL, and batch execution, but not a Spark/Flink replacement

Conclusion

- MLflow provides comprehensive framework for MLOps including performance tracking, model packaging, model serving

-

Conclusion

- Microframework is one of simplest approach to create a microservice
- FastAPI becomes very popular to ML service deployment as it can integrate with other python packages with ease
- FastAPI facilitates input data validation / deserialization, return data serialization, and automatic API document generation
- Combine with docker, ML microservice will be scalable and simple to deploy

References

- P.S. Tan, “Differences: DevOps, ITOps, MLOps, DataOps, ModelOps, AIOps, SecOps, DevSecOps”, <https://medium.com/vitrox-publication/differences-devops-itops-mlops-dataops-modelops-aiops-secops-devsecops-part-1-3-8b238cf72942>
- D. Biswas, “Bridging DataOps and MLOps: ML model inferences as a new Data Source”, <https://towardsdatascience.com/bridging-dataops-mlops-301f010caf30>
- run.ai, “MLflow: the Basics and a Quick Tutorial”, <https://www.run.ai/guides/machine-learning-operations/mlflow>
- T. Danka, “You Should Start Using FastAPI Now”, <https://towardsdatascience.com/you-should-start-using-fastapi-now-7efb280fec02>
- T. Danka, “How to properly ship and deploy your machine learning model”, <https://towardsdatascience.com/how-to-properly-ship-and-deploy-your-machine-learning-model-8a8664b763c4>